

```
In [230]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from scipy import stats, io
from scipy.special import expit
```

```
In [231]: class LogisticRegressionModel(object):

    def __init__(self, name, seed = 42):
        if name != 'data':
            raise ValueError('Incorrect Dataset')
        np.random.seed(seed)
        self.name = name
        self.data = None
        self.trainOrig = None
        self.trLabelsOrig = None
        self.train = None
        self.trLabels = None
        self.val = None
        self.valLabels = None
        self.test = None
        self.dim = None
        self.weight = None
        self.pred = None

        self.load_data()
        self.split(1000)

    def load_data(self):
        self.data = io.loadmat(self.name + '.mat')
        print('Loaded: ' + self.name)
        self.normalize()
        self.trainOrig = self.data['X']
        self.trLabelsOrig = self.data['y'].reshape(-1)
        self.test = self.data['X_test']
        self.test = np.apply_along_axis(self.add1, 1, self.test)

        print('Training size (Before Split): ' + str(len(self.trainOrig)))
        print('Training labels (Before Split):' + str(len(self.trLabelsOrig)))
        print('Test size: ' + str(len(self.test)))

    def split(self, valSize):
        totalLen = len(self.trainOrig)
        trainSize = totalLen - valSize
```

```

randIdx = np.random.permutation(totalLen)
self.train = self.trainOrig[randIdx][:trainSize]
self.train = np.apply_along_axis(self.add1, 1, self.train)
self.trLabels = self.trLabelsOrig[randIdx][:trainSize]
self.val = self.trainOrig[randIdx][trainSize:]
self.val = np.apply_along_axis(self.add1, 1, self.val)
self.valLabels = self.trLabelsOrig[randIdx][trainSize:]
self.dim = self.train.shape

print('Training Data Len: ' + str(len(self.train)))
print('Training Labels Len: ' + str(len(self.trLabels)))
print('Validation Data Len: ' + str(len(self.val)))
print('Validation Labels Len: ' + str(len(self.valLabels)))

def normalize(self):
    for idx in ['X', 'X_test']:
        mean = np.mean(self.data[idx], axis = 0)
        std = np.std(self.data[idx], axis = 0)
        z_score = lambda val: (val-mean)/std
        self.data[idx] = np.apply_along_axis(z_score, 1, self.data[idx])

def add1(self, arr):
    return np.append(arr, [[1]])

def cost(self, X, y, w, l):
    s = expit(X @ w)
    s = np.maximum(s, 1e-7)
    s = np.minimum(s, 1-1e-7)
    return l*w@w - (1/len(X))*(y.dot(np.log(s)) + (1-y).dot(np.log(1-s)))

def grad_penalty(self, X, y, w, l):
    s = expit(X @ w)
    return -X.T.dot(y - s) + 2*l*w

def gradient_descent(self, typeGrad, e, l, changingE, tol=1e-7, graph=True, numIter=1000, maxLoops = 5):
    costLst, valCost = [], []
    converged = False
    w = np.zeros(self.dim[1])
    J_o, J_new = float("Inf"), self.cost(self.train, self.trLabels, w, l)
    i = 0
    if typeGrad == 'batch':
        while i < numIter and not converged:
            w = w - e * self.grad_penalty(self.train, self.trLabels, w, l)
            J_o, J_new = J_new, self.cost(self.train, self.trLabels, w, l)

```

```

        costLst.append(J_o)
        valCost.append(self.cost(self.val, self.valLabels, w,
1))

        i+=1
        if abs(J_o - J_new) < tol:
            converged = True
            break
    if typeGrad == 'stochastic':
        numLoops = 0
        while numLoops < maxLoops and not converged:
            numLoops += 1
            randIdx = np.random.permutation(self.dim[0])
            X = self.train[randIdx]
            y = self.trLabels[randIdx]
            for x_i, y_i in zip(X, y):
                if changingE:
                    w = w - e/(.03*(i+10)) * self.grad_penalty(x_i
, y_i, w, 1)

                else:
                    w = w - e*self.grad_penalty(x_i, y_i, w, 1)
                    J_o, J_new = J_new, self.cost(X, y, w, 1)
                    costLst.append(J_o)
                    valCost.append(self.cost(self.val, self.valLabels,
w, 1))

            i+=1
            if abs(J_o - J_new) < tol:
                converged = True
                break
    print("Cost Final: " + str(J_o))
    print("Converged:" + str(i) + "iterations")
    if graph:
        plt.plot(costLst, label = "Training")
        plt.plot(valCost, label = "Validation")
        plt.xlabel("Number of Iterations")
        plt.ylabel("Cost")
        plt.title("Cost Over Iterations")
    self.weight = w


def experiment(self, typeGrad, testData, e, l, changingE):
    self.gradient_descent(typeGrad, e, l, changingE, graph=False)
    y_val = expit(testData @ self.weight)
    y_pred = (y_val > 0.5).astype(np.int)
    self.pred = y_pred
    return y_pred

def accuracy(self, labels):
    return np.sum(self.pred.reshape(-1) == labels.reshape(-1))/len
(self.pred)

```

```
In [238]: m = LogisticRegressionModel('data', 42)
m.gradient_descent('batch', .01, 1e-6, changingE = False)
```

```
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.03451128789693964
Converged:208iterations
```



```
In [234]: lrs = [10**(-7), 10**(-6), 10**(-5), 10**(-4), 10**(-3), 10**(-2)]
regs = [.01,.001,.0001,.00001,.000001,.0000001]
for lr in lrs:
    for reg in regs:
        mod = LogisticRegressionModel('data', 42)
        mod.experiment('batch', mod.val, lr, reg, changingE = False)
        print("LR: "+str(lr)+' RegParam: '+str(reg)+ ' Accuracy: '+
str(mod.accuracy(mod.valLabels)))
```

```
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
```

Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.5129414392036031  
Converged:1000iterations  
LR: 1e-07 RegParam: 0.01 Accuracy: 0.962  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.5121343753396624  
Converged:1000iterations  
LR: 1e-07 RegParam: 0.001 Accuracy: 0.962  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.5120536688016132  
Converged:1000iterations  
LR: 1e-07 RegParam: 0.0001 Accuracy: 0.962  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.5120455981462916  
Converged:1000iterations  
LR: 1e-07 RegParam: 1e-05 Accuracy: 0.962  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.5120447910807444  
Converged:1000iterations  
LR: 1e-07 RegParam: 1e-06 Accuracy: 0.962  
Loaded: data

```
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.5120447103741896
Converged:1000iterations
LR: 1e-07 RegParam: 1e-07 Accuracy: 0.962
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.2038027086929647
Converged:1000iterations
LR: 1e-06 RegParam: 0.01 Accuracy: 0.977
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.18553734429978758
Converged:1000iterations
LR: 1e-06 RegParam: 0.001 Accuracy: 0.977
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.18371077663228988
Converged:1000iterations
LR: 1e-06 RegParam: 0.0001 Accuracy: 0.977
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
```

Validation Labels Len: 1000  
Cost Final: 0.18352811955325443  
Converged:1000iterations  
LR: 1e-06 RegParam: 1e-05 Accuracy: 0.977  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.1835098538422282  
Converged:1000iterations  
LR: 1e-06 RegParam: 1e-06 Accuracy: 0.977  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.1835080272710942  
Converged:1000iterations  
LR: 1e-06 RegParam: 1e-07 Accuracy: 0.977  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.17324801893735775  
Converged:1000iterations  
LR: 1e-05 RegParam: 0.01 Accuracy: 0.991  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.07048163099075341  
Converged:1000iterations  
LR: 1e-05 RegParam: 0.001 Accuracy: 0.991  
Loaded: data  
Training size (Before Split): 6000

```
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.06020321401375592
Converged:1000iterations
LR: 1e-05 RegParam: 0.0001 Accuracy: 0.991
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.059175354531653845
Converged:1000iterations
LR: 1e-05 RegParam: 1e-05 Accuracy: 0.991
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.05907256840559691
Converged:1000iterations
LR: 1e-05 RegParam: 1e-06 Accuracy: 0.991
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.05906228979121318
Converged:1000iterations
LR: 1e-05 RegParam: 1e-07 Accuracy: 0.991
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
```



Cost Final: 0.357604435955906  
Converged:1000iterations  
LR: 0.0001 RegParam: 0.01 Accuracy: 0.995  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.06444357566896411  
Converged:268iterations  
LR: 0.0001 RegParam: 0.001 Accuracy: 0.992  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.041655862093800636  
Converged:1000iterations  
LR: 0.0001 RegParam: 0.0001 Accuracy: 0.995  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.03877904273626833  
Converged:1000iterations  
LR: 0.0001 RegParam: 1e-05 Accuracy: 0.995  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.038491356274646026  
Converged:1000iterations  
LR: 0.0001 RegParam: 1e-06 Accuracy: 0.995  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000

```
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.0384625875832243
Converged:1000iterations
LR: 0.0001 RegParam: 1e-07 Accuracy: 0.995
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.6568052834526016
Converged:1000iterations
LR: 0.001 RegParam: 0.01 Accuracy: 0.997
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.0973664924757946
Converged:1000iterations
LR: 0.001 RegParam: 0.001 Accuracy: 0.997
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.04034624649051768
Converged:448iterations
LR: 0.001 RegParam: 0.0001 Accuracy: 0.995
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.035192866625506684
```

```
Converged:1000iterations
LR: 0.001  RegParam: 1e-05  Accuracy: 0.997
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.03462706691031586
Converged:1000iterations
LR: 0.001  RegParam: 1e-06  Accuracy: 0.997
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.03457048635840218
Converged:1000iterations
LR: 0.001  RegParam: 1e-07  Accuracy: 0.997
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.8000429683031589
Converged:963iterations
LR: 0.01  RegParam: 0.01  Accuracy: 0.997
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.11250010130894003
Converged:809iterations
LR: 0.01  RegParam: 0.001  Accuracy: 0.997
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
```

Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.04219741338341202  
Converged:625iterations  
LR: 0.01 RegParam: 0.0001 Accuracy: 0.997  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.03534560097392368  
Converged:172iterations  
LR: 0.01 RegParam: 1e-05 Accuracy: 0.996  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.03451128789693964  
Converged:208iterations  
LR: 0.01 RegParam: 1e-06 Accuracy: 0.997  
Loaded: data  
Training size (Before Split): 6000  
Training labels (Before Split):6000  
Test size: 497  
Training Data Len: 5000  
Training Labels Len: 5000  
Validation Data Len: 1000  
Validation Labels Len: 1000  
Cost Final: 0.03442628911709663  
Converged:202iterations  
LR: 0.01 RegParam: 1e-07 Accuracy: 0.997

```
In [240]: #Cell for submission
m = LogisticRegressionModel('data', 42)
y_pred = m.experiment('batch', m.test, .01, 1e-6, changingE = False).r
eshape(-1)

def results_to_csv(y_test):
    y_test = y_test.astype(int)
    df = pd.DataFrame({'Category': y_test})
    df.index += 1 # Ensures that the index starts at 1.
    df.to_csv('submission.csv', index_label='Id')

results_to_csv(y_pred)
```

```
Loaded: data
Training size (Before Split): 6000
Training labels (Before Split):6000
Test size: 497
Training Data Len: 5000
Training Labels Len: 5000
Validation Data Len: 1000
Validation Labels Len: 1000
Cost Final: 0.03451128789693964
Converged:208iterations
```

```
In [ ]:
```